

Project Leader - Gabriel Maliakal ( CMSE )

### **Parallel Tricubic Interpolation to resize 3D Images**

Interpolation of data is very important in all fields of research. It is the process of fitting a curve on ALL points in the data in order to interpolate values in between any two data points that are not actually present in the data. Image interpolation is a prerequisite step in a wide variety of Computer Vision related projects.

Examples of use -

A specific example of this is that I have observed is that during the training of a 3D segmentation network, we have to resize volumes so that they have the same voxel spacing ( also known as isometric images ) in order to make the neural network not have to figure out zoom invariance. I have observed that this is a very important step in 3D medical image segmentation, otherwise the neural network will have poor segmentation as it would be trained on a dataset having varying slice thicknesses.

Packages like `scipy.ndimage` can be used but this is very slow for 3D images. Especially if you want to resize 1000s 3D images to some isometric spacing of  $[1,1,1]$  this takes a very long time. I want to see if I can come up with a faster implementation and call this function in my python scripts.

#### **Serial implementation -**

Before I move onto parallelization, let us look at how I think this works sequentially -

- 1) First step is to compute the size of the new matrix using the zoom factor. If we have a 3D image of size  $Z \times R \times C$ , and if the zoom factor for each dimension is say  $(z,r,c)$  where  $z,r,c > 0$ , then the new size of the resized matrix will be  $zZ \times rR \times cC$ .
- 2) Mapping existing pixels to new locations - The coordinates of the existing pixels will also be multiplied by the zoom factor. For example if there is a pixel at some coordinate  $(d,e,f)$  in the original image, then it will now be in coordinate  $(\text{int}(zd),\text{int}(re),\text{int}(cf))$ . We have to make sure the new coordinates are integers, but this can be done by typecasting.
- 3) Perform 1D cubic interpolation along each axes-
  - a) 1D cubic interpolation for every row voxel, i.e. keep the  $z$  index constant and row index constant, but vary column index from 0 to  $C$ , we get an array of voxels. Do 1D cubic on this for every value of  $z$  index and row index.
  - b) Using the above sub-interpolated matrix, Repeat while keeping the  $C$  and  $Z$  index constant, vary row from 0 to  $R$ , 1D cubic interpolation this array of voxels.
  - c) Using a sub-interpolated matrix obtained from steps a,b now do 1D cubic interpolation while keeping the  $r$  and  $c$  fixed and vary the  $z$  from 0 to  $Z$  for all values of  $r$  and  $c$ .

The process of getting the coefficients of each spline between every consecutive pairs of points is obtained by solving the below equation -

$$Ax = b$$

where  $A$  is the system matrix containing information about the part of the coefficients of each spline and some information using the pairs of points. The vector  $x$  contains the unknown coefficients to be determined while  $b$  contains the linear slope between each consecutive point. The formula for obtaining each component is outlined in [1]. The system matrix  $A$  is a symmetric toeplitz tridiagonal matrix, the equation can be solved using a modified version of Gauss elimination in  $O(n)$  time complexity [2] using Thomas Algorithm.

### **Parallelization strategy -**

- 1) Parallelize the mapping of existing pixels to new locations. This is an embarrassingly parallel operation.
- 2) The next step could be thought of as having two tiers of parallelism (nested) -
  - a) Tier 1 - The process in each sub-step in Step 3) is an embarrassingly parallel process. However, steps a,b,c have to be done sequentially but it could be done in any order, but the steps within step a or b or c are embarrassingly parallel.
  - b) Tier 2 - We can go a step further and try to parallelize the 1D cubic interpolation itself. Instead of using Thomas algorithm which is done in  $O(n)$  complexity we can parallelize the solving of tridiagonal matrix equations using many methods. One particular method is in [3] that involves splitting the tri diagonal matrix into subparts as well as the x vector that is to be solved and the b vector mentioned in the previous section.

**Note** - If I am unable to perform tier 2 then I might just have to use a library like Eigen [4] that performs tridiagonal solving, since this itself can be a major homework/project.

The parallelization can be achieved by native C++ threading and openMP.

Since most operations are embarrassingly parallel a very naive native threading implementation is possible without worrying about mutexes given that I do not perform tier 2 parallelization.

To enable tier 2 parallelism I will have to figure out using multiple mutexes or I can just do it using openMP and let the package figure that out.

I suspect there will be memory issues for very large resizings, so the operations may have to be done using CUDA instead of openMP. I will have to explore how to do it in CUDA.

### **Benchmarking**

To compare my implementation I will be comparing with the scipy and pytorch-

- 1) **Native threading implementation(without tier 2) vs scipy.ndimage.zoom vs torch.nn.functional.interpolate**
- 2) **OpenMP implementation vs scipy.ndimage.zoom vs torch.nn.functional.interpolate**
- 3) **OpenMP implementation Cuda implementation vs scipy.ndimage.zoom vs torch.nn.functional.interpolate**

In the above, I will attempt the first 4 parts with high priority. If time allows I will attempt to improve the algorithm using CUDA, based on the expectation that GPUs handle big data better parallelly than threaded CPUs.

The c++ function will be called in python. This can be done using Boost[5] The numpy array containing the image will be saved as a .npy file and passed to the boost c++ function. The internal c++ function will read this numpy array using another library called cnpy[6].

Strong scaling and weak scaling will be observed for my implementation vs existing implementation.

### **References-**

- [1] [https://www.giassa.net/?page\\_id=274](https://www.giassa.net/?page_id=274)
- [2] <https://arxiv.org/pdf/1402.5094.pdf>
- [3] <https://web.alcf.anl.gov/~zippy/publications/partrid/partrid.html>
- [4] [http://eigen.tuxfamily.org/index.php?title=Main\\_Page](http://eigen.tuxfamily.org/index.php?title=Main_Page)
- [5] [https://www.boost.org/doc/libs/1\\_59\\_0/libs/python/doc/tutorial/doc/html/index.html](https://www.boost.org/doc/libs/1_59_0/libs/python/doc/tutorial/doc/html/index.html)
- [6] <https://github.com/rogersce/cnpy>